

UNIT 6

Concurrency Control

Content

- Types of locks and system lock tables
- Serializability by Two-Phase Locking
- Dealing with Deadlock and Starvation
- Timestamp ordering
- Validation concurrency control

Needs of Lock

Locking technique is used to **control concurrent execution** of transactions.

Enforce **Isolation execution**.

Locks are used as a means of synchronizing the access by concurrent transactions.

Lock:

Lock is a variable associated with a data item that describes the **possible operations that are allowed on data item**.

Types of Lock

- Binary Lock :
- Shared/Exclusive (Read/Write) Lock

Binary Lock

- A Binary Lock can have two state or values:
 - Lock
 - Unlock
- 1 -> indicates locked
- 0 -> indicates unlocked.

Locking

Locking is an operation which secures

- Permission to Read a data item for a transaction.

OR

- Permission to Write a data item for a transaction.

A distinct lock is associated with each database item x.

Example: Lock (X)

Data item X is locked in behalf of the requesting transaction.

Unlocking

Unlocking is an operation which removes these permissions from the data item

Example: Unlock (X).

Data item X is made available to all other transactions.

Exclusive/Shared Lock

Shared Mode: Read lock (X).

More than one transaction can apply share lock on X for reading.

Exclusive mode : Write lock (X).

Only one write lock on X can exist at any time and no shared lock can be applied by any other transaction on X.

If a transaction have exclusive lock then no transaction have exclusive lock on same item.

Lock table and Lock manager

- **Lock table :**

- In the Lock table, system maintain only the items that are concurrently.
- Lock table is organize as a hash file.
- Item not in lock table are consider to be unlocked.

Transaction ID	Data item id	lock mode	Ptr to next data item
T1	X1	Read	Next

- **Lock Manager :**

- Managing locks on data items.

Well-formed Transaction

- Database requires that all transactions should be well-formed.
- A transaction is **well-formed** if:
 1. It must lock the data item before it reads or writes to it.
 2. It must not lock an already locked data items and it must not try to unlock a free data item.

Lock Operation of Binary Lock

Lock_item(X) :

B:

if LOCK (X) = 0 (* item is unlocked*)

then LOCK (X) \leftarrow 1 (*lock the item*)

else

begin

wait (until lock (X) = 0)

the lock manager wakes up the transaction;

goto B

end;

Unlock Operation of Binary Lock

Unlock_item(X) :

LOCK (X) \leftarrow 0 (*unlock the item*)

if (any transactions are waiting) then

 wake up one of the waiting the transactions;

Rules for transaction in binary locking scheme

1. A transaction T must issue the operation `lock_item(X)` before any `read_item(X)` or `write_item(X)` operation performed in T.
2. A transaction T must issue the operation `unlock_item(X)` after all `read_item(X)` and `write_item(X)` operation completed in T.
3. A transaction T will not issue a `lock_item(X)` operation if it already hold lock on item X.
4. A transaction T will not issue a `unlock_item(X)` operation unless it already hold lock on item X.

Lock Operation of Shared Lock

read_item(X) :

B:

if LOCK (X) = “unlocked” then

begin

LOCK (X) \leftarrow “read-locked”;

no_of_reads (X) \leftarrow 1;

end

else if LOCK (X) \leftarrow “read-locked” then

no_of_reads (X) \leftarrow no_of_reads (X) +1

else

begin

wait (until LOCK (X) = “unlocked”

and the lock manager wakes up the transaction);

go to B

end;

Lock Operation of exclusive Lock

write_item(X) :

B:

if LOCK (X) = “unlocked” then

 LOCK (X) \leftarrow “write-locked”;

else

 begin

 wait (until LOCK (X) = “unlocked”

 and the lock manager wakes up the transaction);

 go to B

 end;

Unlock Operation for two-mode

Unlock(x) :

```
if LOCK (X) = "write-locked" then
    begin
        LOCK (X) ← "unlocked";
        wake up on of the waiting transaction, if any
    end
else if LOCK (X) = "read-locked" then
    begin
        no_of_read(X) ← no_of_read(X)-1 ;
        if no_of_read(X) =0 then
            begin
                LOCK (X) ← "unlocked";
                wake up on of the waiting transaction, if any
            end
        end
    end
end;
```


Rules for transaction in shared/exclusive locking scheme

1. A transaction T must issue the operation `read_lock(X)` or `write_lock(X)` before any `read_item(X)` operation performed in T.
2. A transaction T must issue the operation `write_lock(X)` before `write_item(X)` operation performed in T.
3. A transaction T must issue the operation `unlock(X)` after all `write_item(X)` operation are completed in T.

Rules for transaction in shared/exclusive locking scheme

4. A transaction T will not issue a `read_lock(X)` operation if it already holds a `read(shared)` lock or a `write(exclusive)` lock on item X.
5. A transaction T will not issue a `write_lock(X)` operation if it already holds a `read(shared)` lock or a `write(exclusive)` lock on item X.
6. A transaction T will not issue an `unlock(x)` operation unless it already holds a `read(shared)` lock or a `write(exclusive)` lock on item X.

Serializability by Two Phase Locking

- A transaction is said to follow the **two-phase locking protocol** if all locking operations (read_lock, write_lock) precede the first unlock operation in the transaction.
- Such transaction can be divided into 2 phases
 - Expanding or Growing (first) Phase
 - Shrinking (Second) Phase

Serializability by Two Phase Locking

- ❑ **Locking (Growing) Phase:** In this phase new locks on the items can be acquired but none can be released.
 - ❑ A transaction applies locks (read or write) on desired data items one at a time.(Gain Lock)
- ❑ **Unlocking (Shrinking) Phase:** In this phase existing lock can be released but no new lock can be acquired.
 - ❑ A transaction unlocks its locked data items one at a time.(Release Lock)
- **Requirement:** For a transaction these two phases must be mutually exclusively, that is,
 - During locking phase unlocking phase must not start
 - During unlocking phase locking phase must not begin.

Serial Schedule of T1 and T2

<u>T1</u>	<u>T2</u>	<u>Result</u>
read_lock (Y);	read_lock (X);	Initial values: X=20; Y=30
read_item (Y);	read_item (X);	Result of serial execution
unlock (Y);	unlock (X);	T1 followed by T2
write_lock (X);	Write_lock (Y);	X=50, Y=80.
read_item (X);	read_item (Y);	Result of serial execution
X:=X+Y;	Y:=X+Y;	T2 followed by T1
write_item (X);	write_item (Y);	X=70, Y=50
unlock (X);	unlock (Y);	

Problem to violated to phase locking Protocol

T1	T2
<pre>Read_lock(Y) read_item (Y); unlock (Y);</pre>	<pre>read_lock (X); read_item (X); unlock (X); write_lock (Y); read_item (Y); Y:=X+Y; write_item (Y); unlock (Y);</pre>
<pre>write_lock (X); read_item (X); X:=X+Y; write_item (X); unlock (X)</pre>	

Result :

X=50; Y=50

Non-Serializable
because it violated
two phase policy.

T1	T2
Read_lock(Y)	read_lock (X);
read_item (Y);	read_item (X);
write_lock (X);	write_lock (Y);
unlock (Y);	unlock (X);
read_item (X);	read_item (Y);
X:=X+Y;	Y:=X+Y;
write_item (X);	write_item (Y);
unlock (X)	unlock (Y);

- Follows Two phase locking protocols

Two phase locking protocols

- **Advantage:**

- It produce Serializable schedule.

- **Problem With two phase locking protocols**

- Two phase locking protocols can produce deadlock.

Lock conversion

❑ **Lock upgrade:** existing read lock to write lock

if T_i has a read-lock (X) and T_j has no read-lock (X) ($i \neq j$) **then**

convert read-lock (X) to write-lock (X)

else

force T_i to wait until T_j unlocks X

❑ **Lock downgrade:** existing write lock to read lock

T_i has a write-lock (X) (*no transaction can have any lock on X*)

convert write-lock (X) to read-lock (X)

Variation of two phase locking

- Two-phase policy generates two locking algorithms
 - Basic
 - Conservative.
- **Conservative:**
 - Prevents deadlock by locking all desired data items before transaction begins execution by **pre-declaring** its read-set and write-set.
 - Deadlock free protocol
 - If any of pre-declared items can not be locked, then the transaction does not lock any of its items, and it wait until all the item available for locking.

Terms in two-phase locking protocols

☐ Strict 2PL :

- ☐ A transaction must release all its exclusive lock only when it is committed or abort.
- ☐ Transaction T does not release any of its exclusive lock until after commit or aborts.

☐ Rigorous 2PL :

- ☐ A transaction must release all its lock only when it's commit or abort.
- ☐ Transaction T does not release any of its lock(exclusive or Shared) until after commit or aborts.
- ☐ It is easier to implement than strict 2PL.

Difference between Conservative and Rigorous 2PL

- Conservative: Growing phase
- Rigorous : Expanding Phase

Deadlock

- **Deadlock** occurs when each transaction T in a set of two or more transaction is waiting for some item that is locked by some other transaction T' in the set.
- Hence each transaction in the set is on a waiting queue, waiting for one of the other transaction in the set to release the lock on an item.

Deadlock

T1	T2
read_lock (Y) read_item(Y)	read_lock (X) read_item(X)
write_lock (X) (waits for X)	write_lock (Y); (waits for Y)

T1 and T2 did follow two-phase policy but they are deadlock

3. Dealing with Deadlock

- There are several ways to dealing with deadlock
 - Deadlock prevention
 - Deadlock detection
 - Prevent Starvation

Deadlock prevention

- *Deadlock prevention* protocols ensure that the system will *never* enter into a deadlock state.
- Some prevention strategies :
 1. Require that each transaction locks all its data items before it begins execution (pre-declaration).
 - The conservative two-phase locking uses this approach
 2. On the basis of transaction timestamp $TS(T)$

Transaction Timestamp TP(S)

- Transaction timestamp is a **unique identifier assigned to each transaction.**
- It is based on the **order in which transaction** are started.
- If transaction T_1 starts before transaction T_2 then
$$TS(T_1) < TS(T_2)$$
- Older transaction has smaller timestamp value.
- Two scheme that prevent deadlock
 1. Wait-Die
 2. Wound-Wait

Deadlock prevention Scheme.

- **wait-die scheme** — non-preemptive

If $TS(i) < TS(j)$, then (T_i older than T_j)

T_i allowed to wait;

otherwise (T_i younger than T_j)

abort T_i (T_i dies) and

restart it later with the same timestamp.

- Older transaction may wait for younger one to release data item. Younger transactions never wait for older ones; they are rolled back instead.
- A transaction may die several times before acquiring needed data item

- **wound-wait scheme** — preemptive

If $TS(i) < TS(j)$, then (T_i older than T_j)

abort T_j (T_i wounds T_j)

restart it later with the same timestamp.

otherwise (T_i younger than T_j)

T_i allowed to wait

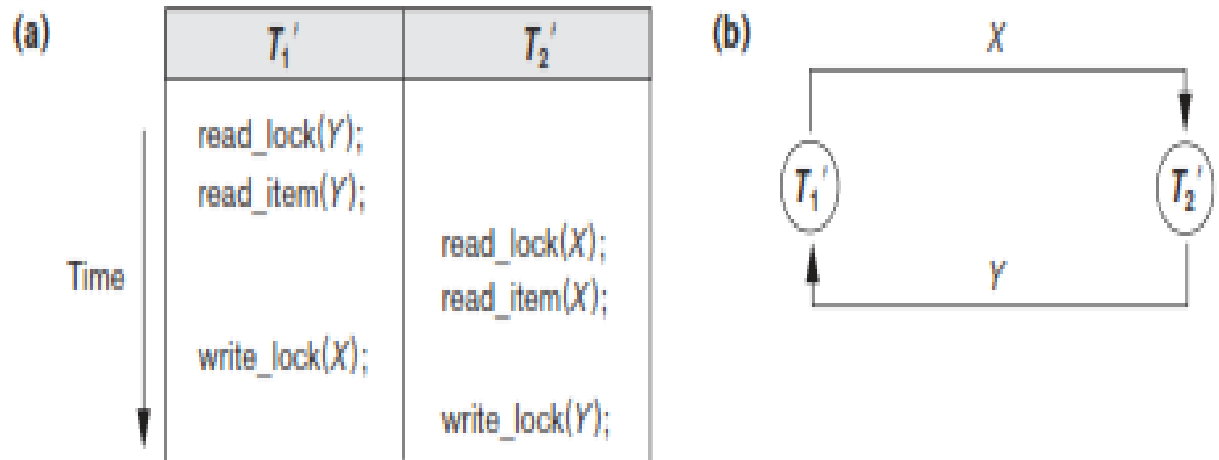
- older transaction *wounds* (forces rollback) of younger transaction instead of waiting for it. Younger transactions may wait for older ones.
- may be fewer rollbacks than *wait-die* scheme.

Deadlock detection and resolution

- In this approach, deadlocks are allowed to happen.
- The scheduler maintains a **wait-for-graph** for detecting cycle.
- If a cycle exists, then one transaction involved in the cycle is selected (victim) and rolled-back.
- A **wait-for-graph** is created using the lock table.
- As soon as a transaction is blocked, it is added to the graph. When a chain like:
 - T_i waits for T_j waits for T_k waits for T_i or T_j occurs, then this creates a cycle.
- One of the transaction of the cycle is selected and rolled back.

wait-for graph

- One node is created in the wait-for graph for each transaction that is currently executing.
- Whenever transaction T_i waiting to lock an item X , that is currently lock by transaction T_j .
- Direct edge is created from $(T_i \longrightarrow T_j)$
- We have a state of deadlock if and only if the wait-for graph has a cycle.



Check whether the following schedule produce deadlock or not?

T1	T2
lock (A) ; read (A) ; write(A); lock(B);	lock(B); write(B); read(B); lock(A);

Apply two phase locking and then check whether schedule is deadlocked or not?	
T1	T2
read (A) ; read(B); If A=1 then B:=B+1 Write(b)	read(B); Read(A) If B=1 then A:=A+1 Write(A)

Starvation

- Starvation :it occurs when a transaction cannot proceed for an indefinite period of time while other transaction in the system continue normally.
- This may occur if the waiting scheme for locked items is unfair.
- Solution of Starvation :
 - First come first served
 - Based on priority
 - But increased the priority of the transaction the longer it waits.
until it get highest priority.

Concurrency control based on timestamp ordering

Timestamp

- Transaction timestamp is a unique identifier assigned to each transaction.
- It is based on the order in which transaction are started.
 - A larger timestamp value indicates a more recent event or operation.
 - Timestamp based algorithm uses timestamp to serialize the execution of concurrent transactions.

Timestamp ordering algorithm

- **Timestamp Ordering** : A schedule in which the transaction participate is then Serializable, and equivalent serial schedule has the transaction in order of their timestamp value this is called timestamp ordering.
- order of accessed item that does not violate serializablility order, for that algorithm associated with each database item X is used.

- **read_TS(X)**

- The read timestamp of item X; this is largest timestamp among all the timestamp of transaction that have successfully read item X.
- $\text{read_TS}(X) = \text{TS}(T)$
 - Where T is the youngest transaction that has read X successfully.

- **write_TS(X)**

- The write timestamp of item X; this is largest timestamp among all the timestamp of transaction that have successfully written item X.
- $\text{write_TS}(X) = \text{TS}(T)$
 - Where T is the youngest transaction that has written X successfully.

Timestamp based concurrency control algorithm

Basic Timestamp Ordering

1. Transaction T issues a write_item(X) operation:

- a. If $\text{read_TS}(X) > \text{TS}(T)$ or if $\text{write_TS}(X) > \text{TS}(T)$, then

abort and roll-back T and reject the operation.

(younger transaction has already read the data item)

- b. If the condition in part (a) does not exist, then

execute write_item(X) of T and set $\text{write_TS}(X) = \text{TS}(T)$.

Timestamp based concurrency control algorithm

Basic Timestamp Ordering

2. Transaction T issues a `read_item(X)` operation:

a. If $\text{write_TS}(X) > \text{TS}(T)$, then

abort and roll-back T and reject the operation.

(younger transaction has already written to the data item)

b. If $\text{write_TS}(X) \leq \text{TS}(T)$, then

execute `read_item(X)` of T and set `read_TS(X)` to the larger of $\text{TS}(T)$ and the current `read_TS(X)`.

- **Advantage :**

- Basic TO always produce conflict Serializable schedule.
- Deadlock does not occur.

- **Problem with Basic TO**

- Cyclic restart may occur if a transaction is continually aborted and restarted.
- A Modification of Basic TO algorithm is known as **Thomas's write rule.**
 - It does not enforce conflict serializability
 - But it reject fewer write operation by modifying the checks for the write_item(X) operation.

Timestamp based concurrency control algorithm

Thomas's Write Rule

1. If $\text{read_TS}(X) > \text{TS}(T)$ then abort and roll-back T and reject the operation.
2. If $\text{write_TS}(X) > \text{TS}(T)$, then just ignore the write operation and continue execution. This is because the most recent writes counts in case of two consecutive writes.
3. If the conditions given in 1 and 2 above do not occur, then execute $\text{write_item}(X)$ of T and set $\text{write_TS}(X)$ to $\text{TS}(T)$.

Validation concurrency control techniques

- optimistic concurrency control technique (validation or certification technique)
 - No checking is done while the transaction is executing.
 - It follows some scheme.
 - In this scheme , updation in the transaction are not applied directly to the database item until the transaction reaches it end.
 - All updation applied to local copies of data item that are kept for transaction.
 - At the end of transaction, concurrency control phases checks whether any of transaction's updates violates serializability or not.

Three phases of concurrency control protocol

- **Read Phase:**

- A transaction can read values of committed data items from the database.
- However , updation are applied only to local copies of data items.

- **Validation Phase :**

- Checking is performed to ensure that serializability will not be violated if transaction updates are applied to the database.

- **Write Phase**

- If validation phase is successful, the transaction updates are applied to the database;
- Otherwise , the updates are discarded and the transaction is restarted
-

Validation Phase

- This phase for T_i checks that, for each transaction T_j that is either committed or is in its validation phase, one of the following conditions holds:
 1. T_j completes its write phase before T_i starts its read phase.
 2. T_i starts its write phase after T_j completes its write phase, and the `read_set` of T_i has no items in common with the `write_set` of T_j
 3. Both the `read_set` and `write_set` of T_i have no items in common with the `write_set` of T_j , and T_j completes its end phase.